Lec-16, IT 527, 24-25

$$\pi'(s) \triangleq \arg\max_a \; q_\pi(s,a)$$

$$= \arg\max_a \; E\left[R_{t+1} + \gamma V_\pi(S_{t+1}) \,\middle|\, S_t = s, A_t = a\right]$$

$$= \arg\max_a \sum_{s',r} p(s',r|s,a)\left[r + \gamma V_\pi(s')\right]$$

$V_\pi(s) = V_{\pi'}(s) \;\; \forall s \in S$ or $S^+ \Rightarrow$

$$\pi'(s) = \arg\max_a \sum_{s',r} p(s',r|s,a)\left[r + \gamma V_{\pi'}(s')\right]$$

$$= \arg\max_a \; q_{\pi'}(s,a)$$

$$V_{\pi'}(s) = q_{\pi'}(s, \pi'(s)) = \max_a q_{\pi'}(s,a)$$

$\Rightarrow V_{\pi'}(s) = \max_a E\left[R_{t+1} + \gamma V_{\pi'}(S_{t+1}) \mid S_t = s, A_t = a \right]$

this is BOE, implying $\pi'(\cdot)$ as one of the optimal policies & $V_{\pi'}(s)$ as the optimal state-value function

↳ $V_*(s)$

$\Pi_* = \{\pi_1, \pi_2\}$

↳ set of optimal policies

Value iteration :- (PE → PI → PItr)

↳ prediction

Policy itr involves policy evaluation which itself is iterative & requires multiple sweeps through the state space.

Is it possible to truncate PE?

without losing the convergence guarantee of Policy itr

What about if PE is ~~stopped~~ after just one sweep?

(one update of each state)

This is called as Value iteration (VIter) → combination of PI & truncated PE steps

$$V_{k+1}(s) \triangleq \max_a E\left[R_{t+1} + \gamma V_k(S_{t+1}) \mid S_t = s, A_t = a\right] \quad \forall s \in S$$

$\pi_0 \xrightarrow{E} V_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V_{\pi_1} \cdots\cdots \pi_* \rightarrow V_{\pi_*}$

VIter

optimal policy π_*

for any policy π, $V_\pi(s)$ is called the state-value func. if it satisfies BE.

BE

BOE